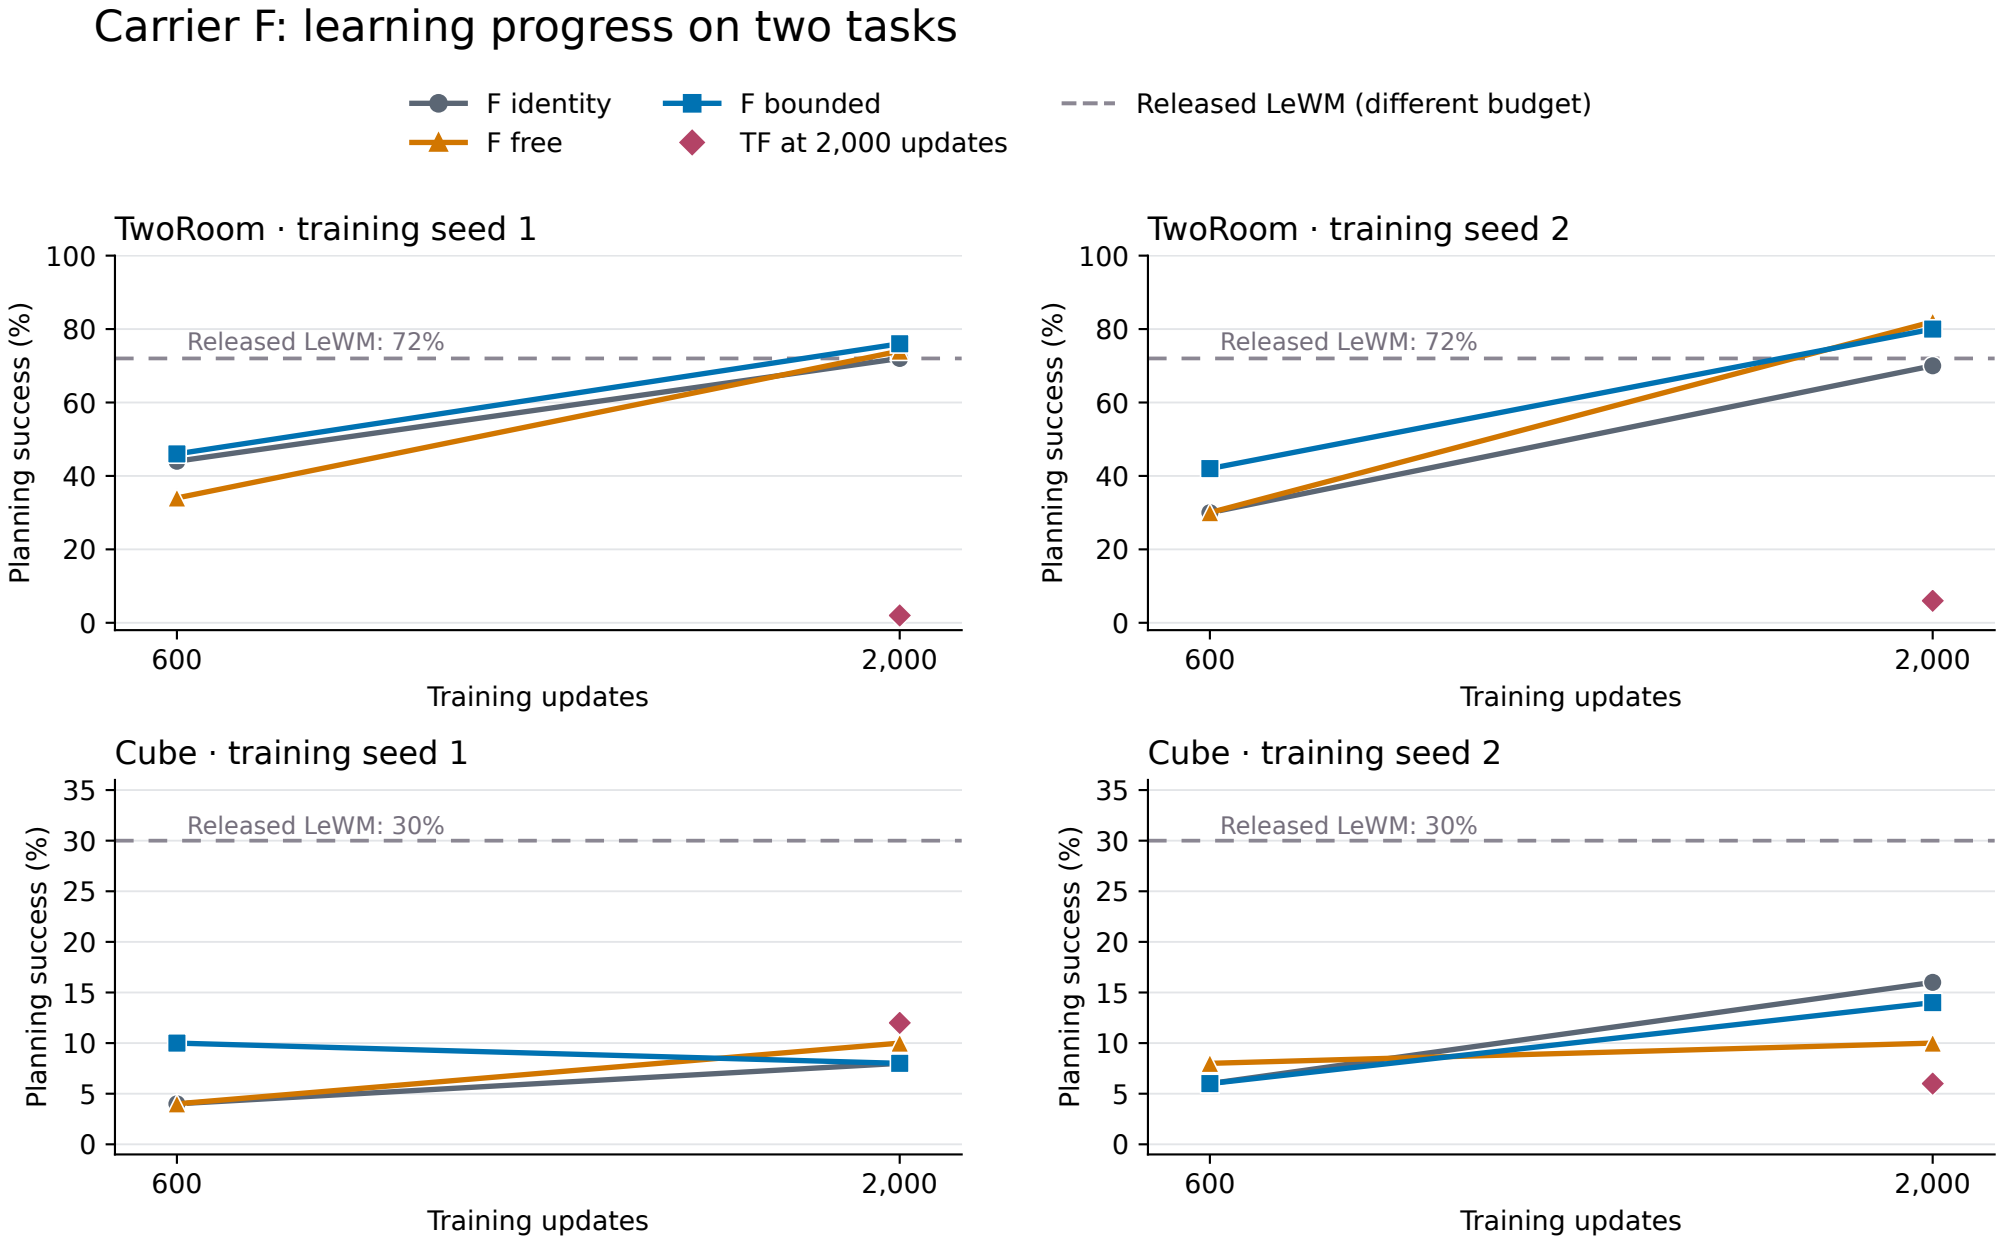

50 fixed validation cases per task; identical planner. Each training seed is shown separately. Rows use different y-scales.
Released LeWM: unmatched training exposure, capacity and normalization. TF: same updates; F changes frozen/trainable capacity.
Segments connect two evaluated checkpoints only; intermediate training performance was not measured.